

Meridian Trade

Turning experimentation into targeted growth

The business problem

Deciding what to pitch runs on guesswork

- We launch crypto, options, and an AI agentic assistant — but which do we show which user?
- Today it runs on hand-maintained rules, so the right offer rarely reaches the right user
- Incentive spend is wasted on users who would never convert
- Experiments aren't measured, so we don't learn what works

The outcome: measured, targeted lift

Product	Control	Treatment	Lift
Agentic assistant	22.7%	39.7%	+17.0 pp
Crypto	26.6%	41.6%	+15.0 pp
Options	23.3%	34.1%	+10.8 pp

One integrated journey on Databricks

- Lakeflow ingests raw activity into a governed medallion in Unity Catalog
- ML propensity models score every user across all three products
- GenAI writes the personalized, compliant outreach nudge
- Lakebase serves it in ~3 ms; Genie answers questions in plain English
- A Databricks app puts it in the Growth team's hands

Built for both buyers

Executive sponsor (VP Growth)

Measured +10 to 17 pp lift

\$441K expected annual value

42% of value in top 30% of users

Domain owner (Growth PM)

Next-best-feature in ~3 ms

Personalized compliant nudges

Self-serve answers via Genie

Proof it's real

Dimension	Result
Scale	10,000 users · 1.55M events
Experiments	30,000 assignments across 3 launches
Models	3 propensity models · AUC 0.65-0.73 · in Unity Catalog
Governance	PII masked · least-privilege grants · lineage captured

What's next

- Close the loop: feed live adoption back as labels and retrain weekly
- Score every future launch against the whole base on day one
- Activate the targeted cohort and nudges into the marketing platform
- Enforce suitability and eligibility guardrails in the recommendation layer



databricks